

Haihao (Sean) Lu

CONTACT INFORMATION	Room 585, Sloan School of Management 100 Main St, Cambridge, MA 02142	Email: haihao@mit.edu Phone: 857-998-3092
RESEARCH INTERESTS	Large-scale and GPU-accelerated optimization; first-order methods; online decision-making with applications in ads; operations research applications.	
APPOINTMENT	Massachusetts Institute of Technology , Cambridge, MA Sloan School of Management The Cecil and Ida Green Career Development Assistant Professor (since July 2025) Assistant Professor of Operations Research and Statistics Google Research , New York City, NY Research scientist (20% part time consulting) University of Chicago , Chicago, IL Booth School of Business Assistant Professor of Operations Management Paternity leave in 2020-2021 Google Research , New York City, NY Visiting Faculty Researcher Large-Scale Optimization Team	September 2024-present December 2022-present July 2020-June 2024 July 2019-June 2020
EDUCATION	Massachusetts Institute of Technology , Cambridge, MA Ph.D. in Mathematics & Operations Research, Thesis Title: Large-Scale Optimization Methods for Data-Science Applications Advisor: Prof. Robert M. Freund Thesis Committee Members: Robert M. Freund, Rahul Mazumder, Ankur Moitra, Jonathan Kelner Shanghai Jiao Tong University , Shanghai, China B.S. in Applied Mathematics, Zhiyuan College, Graduation with distinction Advisor: Prof. David Cai	2014-2019 2010-2014
SELECTED HONORS	INFORMS Computing Society Prize. This annual prize recognizes outstanding papers at the intersection of operations research and computer science. Sloan Research Fellowship (Mathematics). The two-year \$75,000 fellowships are awarded annually to early-career researchers whose creativity, innovation, and research accomplishments make them stand out as the next generation of leaders. COIN-OR Cup winner. The COIN-OR Cup recognizes and celebrates the best contributions to open-source operations research software development and use associated with the COIN-OR software collection.	2026 2026 2024

The Beale–Orchard-Hays Prize. 2024

This prize, awarded every three years by the Mathematical Optimization Society, recognizes excellence in computational mathematical programming.

INFORMS Revenue Management and Pricing Section Prize. 2023

This prize is awarded for the best contribution to the science of pricing and revenue management published in English.

Michael H. Rothkopf Junior Researcher Paper Prize (first place). 2022

This prize recognizes the best paper by a young researcher in the area of Auctions and Market Design.

Special invited session on “spotlight beyond WINE”. 2022

This session at WINE highlights some of the best works in algorithmic game theory that appeared in conferences and journals other than WINE, or mature working papers.

INFORMS Optimization Society Young Researcher Prize. 2021

This prize serves as an esteemed recognition of promising colleagues in the field of optimization who are at the beginning of their academic or industrial career.

RESEARCH GRANTS

GE Vernova 2026-2027

“GPU Solver Engine for Grid Orchestration at Scale.” My share: \$150,000.

Air Force Office of Scientific Research 2024-2027

“Scalable First-Order Methods for Large Scale Optimization: Theory and Computation.” Award: \$400,000.

Office of Naval Research 2024-2027

“Collaborative Proposal: Scaling up Linear Programming with First-Order Methods and MPI.” My share: \$456,409.

MIT Social and Ethical Responsibilities of Computing 2025

“Addressing Regressive Property Taxation with Fairness-Constrained Differentiable Machine Learning.” Award: \$75,000.

MIT Research Support Committee 2025

Award: \$90,000.

MIT Sloan Junior Faculty Research Assistance Program 2025

Award: \$30,000.

MIT Sloan Junior Faculty Research Assistance Program 2024

Award: \$30,000.

eBay Research Gift 2022-2024

“Efficient Algorithms for Solving Large Scale Choice-Based Model and Applications to Recommendation System.” Award: \$100,000.

JOURNAL PUBLICATIONS¹

1. “PDOT: a Practical Primal-Dual Algorithm and a GPU-Based Solver for Optimal Transport”, with Jinwen Yang, to appear in *SIAM Journal on Optimization*.
2. “PDLP: A Practical First-Order Method for Large-Scale Linear Programming”, with David Applegate, Mateo Díaz, Oliver Hinder, Miles Lubin, Brendan O’Donoghue, and Warren Schudy, *Mathematical Programming Computation* (2026).

¹Publications denoted as *with coauthors* follow alphabetical author order per convention in the field; publications with full author lists follow non-alphabetical order reflecting contributions.

3. “Analysis of Dual-Based PID Controllers through Convolutional Mirror Descent”, with Santiago R. Balseiro, Vahab Mirrokni and Balasubramanian Sivan, to appear in *Operations Research*.
4. “A New Crossover Algorithm for LP Inspired by the Spiral Dynamic of PDHG”, with Tianhao Liu, forthcoming in *INFORMS Journal on Computing*.
5. “cuPDL.jl: A GPU Implementation of Restarted Primal-Dual Hybrid Gradient for Linear Programming in Julia”, with Jinwen Yang, *Operations Research*, 73.6 (2025): 3440-3452.
6. “On the Sparsity of Optimal Linear Decision Rules for a Class of Robust Optimization Problems with Box Uncertainty Sets”, with Bradley Sturt, *Operations Research*, 74.1 (2026): 500-516.
7. “A Practical and Optimal First-Order Method for Large-Scale Convex Quadratic Programming”, with Jinwen Yang, *Mathematical Programming*, 215.1 (2026): 771-808.
8. “Optimizing Scalable Targeted Marketing Policies with Constraints”, with Duncan Simester and Yuting Zhu, *Marketing Science*, 44.5 (2025): 1082-1103.
9. “Regularized Online Allocation Problems: Fairness and Beyond”, with Santiago Balseiro and Vahab Mirrokni, *Manufacturing & Service Operations Management*, 27.3 (2025): 720-735.
10. “Auto-bidding and Auctions in Online Advertising: A Survey”, with many people from Google, *ACM SIGecom Exchanges* 22.1 (2024): 159-183.
11. “On the Geometry and Refined Rate of Primal-Dual Hybrid Gradient for Linear Programming”, with Jinwen Yang, *Mathematical Programming*, 212 (2025): 349-387.
12. “Infeasibility Detection with Primal-Dual Hybrid Gradient for Large-Scale Linear Programming”, with David Applegate, Mateo Diaz and Miles Lubin, *SIAM Journal on Optimization* 34.1 (2024): 459-484.
13. “On the Linear Convergence of Extra-Gradient Methods for Nonconvex-Nonconcave Minimax Problems”, Saeed Hajizadeh, Haihao Lu and Benjamin Grimmer, *INFORMS Journal on Optimization* 6.1 (2024): 19-31.
14. “A J-Symmetric Quasi-Newton Method for Minimax Problems”, with Azam Asl and Jinwen Yang, *Mathematical Programming* 204.1 (2024): 207-254.
15. “The landscape of the proximal point method for nonconvex–nonconcave minimax optimization”, Benjamin Grimmer, Haihao Lu, Pratik Worah and Vahab Mirrokni, *Mathematical Programming* 201.1-2 (2023): 373-407.
16. “Faster First-Order Primal-Dual Methods for Linear Programming using Restarts and Sharpness”, with David Applegate, Oliver Hinder and Miles Lubin, *Mathematical Programming* 201.1-2 (2023): 133-184.
17. “Frank-Wolfe Methods with an Unbounded Feasible Region and Applications to Structured Learning”, Haoyue Wang, Haihao Lu and Rahul Mazumder, *SIAM Journal on Optimization* 32.4 (2022): 2938-2968.

18. “The Best of Many Worlds: Dual Mirror Descent for Online Allocation Problems”, with Santiago Balseiro and Vahab Mirrokni, *Operations Research* 71.1 (2022): 101-119.
19. “An $O(s^r)$ -Resolution ODE Framework for Discrete-Time Optimization Algorithms and Applications to the Linear Convergence of Minimax Problems”, Haihao Lu, *Mathematical Programming* 194 (2022): 1061–1112.
20. “Randomized Gradient Boosting Machines”, Haihao Lu and Rahul Mazumder, *SIAM Journal on Optimization* 2020, 30(4), 2780-2808.
21. “Generalized Stochastic Frank-Wolfe Algorithm with Stochastic ‘Substitute’ Gradient for Structured Convex Optimization”, Haihao Lu and Robert M. Freund, *Mathematical Programming* 187.1 (2021): 317-349.
22. “‘Relative-Continuity’ for Non-Lipschitz Non-Smooth Convex Optimization using Stochastic (or Deterministic) Mirror Descent”, Haihao Lu, *INFORMS Journal on Optimization*, 2019, 1(4): 288-303.
23. “Relatively-Smooth Convex Optimization by First-Order Methods, and Applications”, Haihao Lu, Robert M. Freund and Yurii Nesterov, *SIAM Journal on Optimization*, 2018, 28(1): 333-354.
24. “New Computational Guarantees for Solving Convex Optimization Problems with First Order Methods, via a Function Growth Condition Measure”, Robert M. Freund, Haihao Lu, *Mathematical Programming* 2018, Vol.170, No.2: 445-477.
25. “Stochastic linearization of turbulent dynamics of dispersive waves in equilibrium and nonequilibrium state”, Shi-xiao W. Jiang, Haihao Lu, Douglas Zhou, and David Cai, *New Journal of Physics*, 2016, 18(8): 083028.
26. “Renormalized Dispersion Relations of β -Fermi-Pasta-Ulam Chains in Equilibrium and Nonequilibrium states”, Shi-xiao W. Jiang, Haihao Lu, Douglas Zhou, and David Cai. *Physical Review E*, 2014, 90(3): 032925.

REFEREED
CONFERENCE
PUBLICATIONS

1. “A Field Guide for Pacing Budget and ROS Constraints”, with Santiago Balseiro, Kshipra Bhawalkar, Zhe Feng, Vahab Mirrokni, Balasubramanian Sivan and Di Wang, *ICML 2024*.
2. “Online Ad Procurement in Non-stationary Autobidding Worlds”, with Jason Liang and Baoyu Zhou, *NeurIPS 2023*.
3. “Limiting Behaviors of Nonconvex-Nonconcave Minimax Optimization via Continuous-Time Systems”, Benjamin Grimmer, Haihao Lu, Pratik Worah and Vahab Mirrokni, *ALT 2022*.
4. “Practical Large-Scale Linear Programming using Primal-Dual Hybrid Gradient”, with David Applegate, Oliver Hinder, Mateo Diaz, Miles Lubin, Brendan O’Donoghue, and Warren Schudy, *NeurIPS 2021*.
5. “Regularized Online Allocation Problems: Fairness and Beyond”, with Santiago Balseiro and Vahab Mirrokni, *ICML 2021*.
6. “Contextual Reserve Price Optimization in Auctions via Mixed Integer Programming”, Joey Huchette, Haihao Lu, Hossein Esfandiari and Vahab Mirrokni, *NeurIPS*, 2020.

7. “Dual Mirror Descent for Online Allocation Problems”, with Santiago Balseiro and Vahab Mirrokni, *ICML*, 2020.
8. “Accelerating Gradient Boosting Machines”, Haihao Lu, Sai Praneeth Karimireddy, Natalia Ponomareva and Vahab Mirrokni, *AISTATS*, 2020.
9. “Ordered SGD: A Stochastic First-Order Method for Ordered Empirical Risk Minimization”, Kenji Kawaguchi and Haihao Lu, *AISTATS*, 2020.
10. “Accelerating Greedy Coordinate Descent Methods”, Haihao Lu, Robert M. Freund and Vahab Mirrokni, *ICML*, 2018.
11. “Approximate Leave-One-Out for Fast Parameter Tuning in High Dimensions”, Shuaiwen Wang, Wenda Zhou, Haihao Lu, Arian Maleki, Vahab Mirrokni, *ICML*, 2018.

REFEREED
WORKSHOP
PUBLICATIONS

1. “MPAX: Mathematical Programming in JAX”, with Zedong Peng and Jinwen Yang, *NeurIPS Workshop on Scalable Optimization*, 2025.

PREPRINTS AND
PAPERS UNDER
REVIEW

1. “PDNQP: A GPU-based Factorization-free Method for Large-scale Nonconvex Quadratic Programming”, with Zixi Chen.
2. “On the Local Linear Convergence of Operator Splitting Methods for Conic Programming”, with Lijun Ding and Jinwen Yang.
3. “FlashFolio: A GPU-Accelerated Solver for Portfolio Optimization”, with Yilun Jiang, Zedong Peng and Jinwen Yang.
4. “End-to-End Learning of Correlated Operating Reserve Requirements in Security-Constrained Economic Dispatch”, Owen Shen, Hung-po Chao, Haihao Lu and Patrick Jaillet.
5. “Active Set Identification and Rapid Convergence for Degenerate Primal-Dual Problems”, with Mateo Díaz, Pedro Izquierdo Lehmann and Jinwen Yang.
6. “cuPDLpx: A Further Enhanced GPU-Based First-Order Solver for Linear Programming”, with Zedong Peng and Jinwen Yang.
7. “Enhanced PDHG for Linear Programming with Online Preconditioning”, with Wanyu Zhang.
8. “An Overview of GPU-based First-Order Methods for Linear Programming and Extensions”, with Jinwen Yang.
9. “New Understandings and Computation on Augmented Lagrangian Methods for Low-Rank Semidefinite Programming”, with Lijun Ding and Jinwen Yang.
10. “MPAX: Mathematical Programming in JAX”, with Zedong Peng and Jinwen Yang.
11. “Restarted Halpern PDHG for Linear Programming”, with Jinwen Yang.
12. “The Power of Linear Programming in Sponsored Listings Ranking: Evidence from a Large-Scale Field Experiment”, with Luyang Zhang and Yuting Zhu.

13. “cuPDLP-C: A Strengthened Implementation of cuPDLP for Linear Programming by C Language”, Haihao Lu, Jinwen Yang, Haodong Hu, Qi Huangfu, Jinsong Liu, Tianhao Liu, Yinyu Ye, Chuwen Zhang and Dongdong Ge.
14. “Achieving Fairness and Accuracy in Regressive Property Taxation”, with Ozan Candogan and Feiyu Han.
15. “On the Convergence of L-shaped Algorithms for Two-Stage Stochastic Programming”, with John R. Birge and Baoyu Zhou.
16. “On a Unified and Simplified Proof for the Ergodic Convergence Rates of PPM, PDHG and ADMM”, with Jinwen Yang.

TECHNICAL REPORTS

1. “Depth Creates No Bad Local Minima”, Haihao Lu and Kenji Kawaguchi, *Technical Report*.
2. “Nearly Optimal Linear Convergence of Stochastic Primal-Dual Methods for Linear Programming”, with Jinwen Yang.
3. “Approximate Leave-One-Out for High-Dimensional Non-Differentiable Learning Problems”, Shuaiwen Wang, Wenda Zhou, Arian Maleki, Haihao Lu and Vahab Mirrokni.

TEACHING EXPERIENCE

MIT:²

Instructor

- 15.071. The Analytics Edge
 - Two sections (MBA): 6.0/7, 4.3/5. Spring 2026
 - One section (MBA): 6.5/7, 4.4/5. Spring 2025
- 15.C57/15.C571/6.C57/6.C571/IDS.C57. Optimization
 - Two sections (Undergraduate and Master): 5.1/7, 4.0/5. Fall 2024

Guest Instructor

- 6.251/15.081. Introduction to Mathematical Programming
 - Two lectures in Fall 2022
 - Two lectures in Fall 2024

Booth School of Business, University of Chicago:³

Instructor

- BUSN 32100 Data Analysis with R and Python (MBA)
 - Section 32100-01 (Full-Time MBA): 4.6, 4.6, 4.7, 4.6, 4.6. Winter 2024
 - Section 32100-02 (Full-Time MBA): 4.6, 4.4, 4.8, 4.5, 4.5. Winter 2024
 - Section 32100-81 (Evening MBA): 4.2, 3.8, 4.4, 4.4, 3.9. Winter 2024
 - Section 32100-01 (Full-Time MBA): 4.8, 4.7, 4.7, 4.5, 4.7. Spring 2023
 - Section 32100-02 (Full-Time MBA): 4.9, 4.7, 4.8, 4.7, 4.7. Spring 2023
 - Section 32100-01 (Full-Time MBA): 4.7, 4.5, 4.7, 4.3, 4.5. Spring 2022
 - Section 32100-81 (Evening MBA): 4.8, 4.7, 4.9, 4.6, 4.8. Spring 2022

²Two evaluation questions: “Overall rating” (out of 7), “Recommend Professor” (out of 5).

³Five evaluation questions (out of 5): “Did the instructor convey the course material clearly,” “Did the instructor convey the material in an interesting way,” “Did you take away useful tools, concepts, and/or insights from this course,” “How much did you get out of this course,” and “Would you recommend this course to other students.”

- Section 32100-02 (Full-Time MBA): 5.0, 4.3, 5.0, 4.8, 5.0. Spring 2021
- BUSN 36919 Modern Large-Scale Optimization: Theory and Computation (PhD)
 - Section 36919-50: 4.8, 4.5, 5.0, 4.5, 4.5. Spring 2023

STUDENTS

Current Members

Ph.D. Students (Theses in Progress)

- Charlie Liu (MIT ORC), co-advised with Rahul Mazumder, expected 2029
- Jacob Dentes (MIT ORC), expected 2030

Postdoctoral Researchers

- Zedong Peng (MIT)
- Bo Tang (MIT)
- Vinit Ranjan (MIT)
- Yuchen Lou (MIT)

Alumni

Ph.D. Students

- Jinwen Yang (UChicago Statistics)
Next position: Assistant Professor, Columbia University

Postdoctoral Researchers

- Azam Asl (UChicago Booth)
Next position: consulting firm
- Baoyu Zhou (UChicago Booth)
Next position: Assistant Professor, Arizona State University

Master's Theses Supervised

- Feiyu Han (UChicago Statistics), M.Sc. 2023
Next position: Ph.D. program, UChicago Booth
- Luyang Zhang (UChicago Applied Mathematics), M.Sc. 2022
Next position: Ph.D. program, CMU Heinz
- Hanyang Jiang (UChicago Applied Mathematics), M.Sc. 2021
Next position: Ph.D. program, Georgia Tech ISE

Undergraduates

- Luke Fitzgerald (MIT Mathematics)
- Yang Meng (UChicago CS)
Next position: Ph.D. program, UChicago CS
- Wanyu Zhang (Shanghai University of Finance and Economics)
Next position: Ph.D. program, Stanford MS&E
- Hongpei Li (Shanghai University of Finance and Economics)
Next position: Ph.D. program, Northwestern IEMS

Additional Advising

MIT MBAn Capstone Advising

- Yi Wang and Joyce Zhang 2026
- Felicia Ann Nappi and Gabe Miller 2025

Service on Thesis or Dissertation Committees

- Zikai Xiong (MIT ORC), Ph.D. 2025, advised by Robert M. Freund
- Xiang Meng (MIT ORC), Ph.D. expected 2026, advised by Rahul Mazumder
- Junhui Zhang (MIT EECS), Ph.D. expected 2026, advised by Patrick Jaillet

ACADEMIC SERVICE

Editorial Service

- Associate Editor, *Mathematics of Operations Research* 2026-present
- Associate Editor, *Computational Optimization and Applications* 2026-present
- Editorial Board of Reviewers, *Journal of Machine Learning Research* 2020-present

Professional Society Leadership

- Vice Chair (Linear and Conic Optimization), INFORMS Optimization Society 2024-2026

Conference and Workshop Organization

- Stream Organizer, International Symposium on Mathematical Programming (scheduled) 2027
- Program Committee, INFORMS Optimization Society Conference 2026, 2024
- Cluster Chair, INFORMS Annual Meeting 2025, 2024
- Organizer, Workshop on Modern Continuous Optimization at MIT 2023

Award Committees

- Member, Land-Doig MIP Competition Committee 2026
- Chair, COIN-OR Cup Committee 2025
- Judge, INFORMS George Nicholson Student Paper Competition 2024, 2023

University Committee Service

- MIT Social and Ethical Responsibilities of Computing (SERC) Seed Grant Review Committee 2026
- MIT Sloan OR/Statistics Group Hiring Committee 2025
- MIT ORC Master's Program Admissions Committee 2025
- MIT Sloan MBAn Program Admissions Reader 2024, 2025
- Chicago Booth OM Faculty Recruiting Committee Member 2022-2023
- Chicago Booth OM Workshop Co-Organizer 2021-2022

Reviewer for Journals and Conferences:

Mathematical Programming, SIAM Journal on Optimization, SIAM Journal on Control and Optimization, Mathematics of Operations Research, Mathematical Programming Computation, INFORMS Journal on Computing, Foundations of Computational Mathematics, Operations Research, Management Science, Manufacturing & Service Operations Management, Production and Operations Management, Journal of Machine Learning Research, IEEE Transactions on Image Processing, Computational Optimization and Applications, INFORMS Journal on Optimization, Machine Learning, Journal of Optimization Theory and Applications, NeurIPS, WebConf, ICML.

Professional Memberships:

Institute for Operations Research and the Management Sciences (INFORMS); Mathematical Optimization Society (MOS); Society for Industrial and Applied Mathematics (SIAM).

SELECTED PRESENTATIONS

- 2026** (scheduled) University of Chicago; Booth School of Business
 (scheduled) Duke University; Fuqua School of Business
 (scheduled) MIT; CSE Seminar:
 (scheduled) Keynote Talk at Workshop on Optimisation, Metric Bounds, Approximation and Transversality
 (scheduled) Keynote Talk at Australian Mathematical Society Conference 2026

- (scheduled)* Curtin University; Mathematics
- (scheduled)* RMIT University; Mathematics
- (scheduled)* INFORMS Annual Meeting
- Invited speaker at AFOSR Mathematical Optimization Workshop
- Zuse Institute in Berlin
- Invited speaker at workshop on Foundations of Computational Mathematics
- Invited speaker at Mostly OM workshop
- Invited speaker at workshop on Foundation Models of the Electric Grid
- Amazon; STOC team
- 2025** Plenary speaker at Cargese Workshop on Combinatorial Optimization
- INFORMS Annual Meeting
- Carnegie Mellon University; Tepper School of Business
- NeurIPS Workshop on ScaleOPT: GPU-Accelerated and Scalable Optimization
- Invited speaker at AFOSR Mathematical Optimization Workshop
- International Conference on Continuous Optimization (ICCOPT)
- MIP Workshop
- Invited speaker at workshop on Strategies for Handling Applications with Non-convexity
- Invited speaker at workshop on Combinatorial Optimization and Data Science
- Gurobi Inc
- 2024** NVIDIA; cuOpt team
- Cornell University; Operations Research and Information Engineering
- MIT; Operations Research Center
- Harvard University; School of Engineering and Applied Sciences
- INFORMS Annual Meeting
- Invited speaker at AFOSR Mathematical Optimization Workshop
- Google Research; Cambridge
- Stanford University; Workshop for Yinyu Ye's Retirement Celebration
- Northwestern University; Industrial Engineering and Management Sciences
- University of Minnesota; Industrial and Systems Engineering
- International Symposium on Mathematical Programming (ISMP)
- INFORMS Conference on Optimization
- 2023** MIT; Operations Research Center
- Columbia University; DRO and IEOR joint seminar
- University of Southern California; Viterbi School of Engineering
- Chinese University of Hong Kong; School of Data Science
- INFORMS Annual Meeting
- SIAM Conference on Optimization
- Shanghai University of Finance and Economics
- 2022** INFORMS Annual Meeting
- Argonne National Laboratory
- Johns Hopkins University; Applied Mathematics Department
- Cornell University; Operations Research and Information Engineering
- MIT; Operations Research Center
- Tsinghua University; Institute for Interdisciplinary Information Sciences
- Simons Institute for the Theory of Computing
- 2021** University of Chicago; Booth School of Business
- Carnegie Mellon University; Tepper School of Business
- MIT; Institute for Data, Systems, and Society
- New York University; Stern School of Business
- University of Illinois at Chicago; Liautaud Graduate School of Business
- 2020** eBay
- Google Research; New York City
- International Conference on Artificial Intelligence and Statistics

- Chinese Academy of Sciences
 Shanghai Jiao Tong University; AI+Math Colloquia
 INFORMS Annual Meeting
 SIAM Conference on Mathematics of Data Science
2019 International Conference on Continuous Optimization (ICCOPT)
 Google Research; New York City
 Rensselaer Polytechnic Institute; Mathematics Department
 New York University; Courant Institute of Mathematical Sciences
 National University of Singapore; ISEM
 University of Toronto; Mathematics and Computer Science
 University of Toronto; MIE
 University of Chicago; Booth School of Business
 University of Minnesota, Twin Cities; ISE
 Duke University; Mathematics and Computer Science
 Columbia University; IEOR
 University of Wisconsin–Madison; ISE
 University of Illinois Urbana-Champaign; ISE
 Google Research; Cambridge
2018 INFORMS Annual Meeting
 Princeton University; Princeton Optimization Day
 Google Research; New York City
 Columbia University; Statistics Department
 International Symposium on Mathematical Programming (ISMP)
 International Conference on Machine Learning
 NYAS Meeting on Machine Learning
 INFORMS Conference on Optimization
2017 INFORMS Annual Meeting
 SIAM Conference on Optimization
2016 INFORMS Annual Meeting
 University of Edinburgh
2015 International Symposium on Mathematical Programming (ISMP)

MEDIA AND BLOGS

Google research blog: Scaling up linear programming with PDLP ([link](#)).
 SIAG on Optimization Views and News: First-Order Methods for Linear Programming ([link](#)).
 NVIDIA technical blog: Accelerate Large Linear Programming Problems with NVIDIA cuOpt ([link](#)).
 Quantum Zeitgeist: Google’s PDLP Solver Revolutionizes Linear Programming Problem Solving Efficiency ([link](#)).
 Marktechpost: PDLP (Primal-Dual Hybrid Gradient Enhanced for LP): A New FOM–based Linear Programming LP Solver that Significantly Scales Up Linear Programming LP Solving Capabilities ([link](#)).
 TMT Post: Cardinal Optimizer is ushering in a new “Age of Exploration” for operations research with GPU chips (in Chinese, [link](#)).

MISCELLANEOUS

Computing: Python, Julia, C++, R
Hobbies: Food/Cuisine, Kayaking, Sailing, Hiking, Skiing